# Yasir Raza

yasirabdali6@gmail.com | github.com/Yasirrazaa | linkedin.com/in/yasir-raza | Islamaba

+92 314 5363379 | pypi.org/user/yasirabdali6 | kaggle.com/yasirabdaali | yasirrazaa.github.io

Summary: ML Engineer with 1.5+ years of remote experience on US-based teams, owning the full ML lifecycle: data pipelines, training, rigorous evaluation, and GPU deployment. Built the verified training + evaluation stack for an RL (**GRPO**) math LLM; shipped a **4.5×** TTS inference speedup; processed thousands of tools and job postings into production datasets.

## Technical Skills

Modeling & Training: PyTorch, Transformers, scikit-learn, LoRA/QLoRA, Unsloth, PEFT, GRPO (TRL), supervised fine-tuning, reward design

Data Engineering: pandas, NumPy, SQL, Scrapy, Playwright, deduplication at scale, feature pipelines, S3

Evaluation & Stats: Execution-based harnesses, frozen/contamination-checked splits, bootstrap 95% CIs, WER, RTF benchmarks, time-series forecasting

MLOps & Deployment: Docker, FastAPI, CI/CD (GitHub Actions), RunPod/AWS serverless, model serving, GPU cost optimization

Languages: Python, SQL, Bash, TypeScript

## Work Experience

Teamlift — United States (Remote)

**Junior Data Scientist** **Jan 2024 - Dec 2024**

- Scraped **thousands of AI tools** from public tool directories, building the structured dataset that powers **Teamo**, Teamlift's recommendation engine.
- Built an automated **verification pipeline** (HTTP status checks, screenshot capture, freshness validation) that kept listings accurate and trustworthy — eliminating stale and duplicate entries at scale.
- Preprocessed and normalized thousands of tool records into model-ready features for the recommendation engine.
- Worked directly with US-based stakeholders remotely, shipping dataset improvements on tight iteration cycles.

**Data Science Intern** **Jul 2023 - Dec 2023**

- Built an automated pipeline collecting **thousands of job postings** from dice.com with **deduplication** — hardened with retries, rate limits, and robust parsers.
- Extracted skill-demand signals from job descriptions at scale; trained **time-series models** to forecast demand, informing product and market decisions.

## Technical Projects

**Math2Code — LaTeX → Verified Python via Rule-Based RL** **2025 - 2026**

- Built the complete verified training + evaluation stack for a math LLM (**GRPO/RLVR** on TRL): **22,796 deduplicated samples**, SHA-256 frozen split, oracle-verified gold solutions (**397/397**), and a sandbox pool sustaining **389 exec/s**. Total budget under **$100**.

**Kausal — Production AI Agent Platform (private)** **2026**

- Extended the ii-agent framework in **pure Python** (no LangGraph): custom **RAG**, layered memory, **OpenTelemetry** observability (tracing, metrics, SigNoz dashboards), realtime **voice/video via LiveKit**, **LiteLLM** routing, agent skills, S3 storage — guarded by a **7,275-test** suite.

**Chatterbox-Optimized — Production TTS Engine** **2025**

- Up to **4.5× faster inference** (RTF 0.59 → 0.13) via bucketed CUDA Graphs, bfloat16, and KV-cache tuning; Whisper-validated candidate selection and a reproducible benchmarking suite on difficult text corpora.

**NotebookLLM — Token-Efficient Notebook Tooling** **2025 - 2026**

- Open-source package + MCP server (**10K+ PyPI downloads**): token-budget modes drop up to 80% of context waste, tiktoken measurement per cell, and output summarization for DataFrames, plots, and tracebacks.

**Archon 2.0 — Agent Security Platform** **2026**

- 8-layer defense pipeline, **222-probe corpus** across 10 packs, runtime defense proxy, full **OWASP ASI01–10** coverage; AgentDojo v1 benchmark published (0% ASR on direct overrides at the zero-cost deterministic tier) — **2,295 tests**, 93% gate.

**PRISM-Bench — Cultural Intelligence Benchmark** **2026**

- Authored **PRISM: Pluralistic Reasoning & Identity-Specific Modeling**, an open benchmark measuring cultural intelligence (CQ) across AI systems.

**Forge — The API Compiler for the Agentic Era** **2024 - 2026**

- One OpenAPI 3.x spec in, production artifacts out: SDKs in **10 languages**, CLIs, MCP servers, Agent Skills — **17 output types**; TypeScript SDKs compile strict on Stripe ( 3K files) and Cloudflare ( 14K files) specs.

## Education

**B.Sc. Software Engineering**, National University of Modern Languages (NUML), Islamabad **Aug 2026**

## Certifications

**Datasets Expert · Notebooks Expert (13 medals; Datasets rank 731/11,277, best #109)**, Kaggle **2024 - 2026**

**Associate Data Scientist in Python · Data Analyst with Python (Career Tracks)**, DataCamp **2025 - 2026**

**Generative AI Fundamentals · Intro to LLMs · Intro to Generative AI · Intro to Responsible AI**, Google Cloud Skill Badges **2023**

**Python · Intro to Programming (Professional Certificates)**, Kaggle Learn **2025**

**Advanced LLM Bootcamp**, NUST & Sky Electric **2024**

**Claude Certified Developer – Foundations (proctored exam)**, Anthropic **2026**

**Aspire Leaders Program — Cohort 4 2026**, Aspire Institute (founded at Harvard) **2026 (expected)**

in progress **McKinsey.org Forward Program — Core Skills**, McKinsey.org **2026 (expected)**

in progress